

workshop-03-tasks

March 14, 2021

1 Demo

- Data pre-processing
- Nearest neighbors based classification

```
[30]: # Import the commonly-used modules
import pandas as pd
import numpy as np

# Load the dataset
data = pd.read_csv("/home/dalia/Master of Information Technology in_
↳Cybersecurity/Session 3/COMP8325 Artificial Intelligence/week_3/Week 3_
↳Workshop-20210310/data/titanic.csv")
print("data size: "+str(data.shape))
data.head()
```

data size: (891, 12)

```
[30]: PassengerId      Name      Sex \
0      1      Braund, Mr. Owen Harris    male
1      2  Cumings, Mrs. John Bradley (Florence Briggs Th...  female
2      3      Heikkinen, Miss. Laina    female
3      4  Futrelle, Mrs. Jacques Heath (Lily May Peel)  female
4      5      Allen, Mr. William Henry    male

      Age  SibSp  Parch      Ticket    Fare Cabin Embarked  Pclass \
0  22.0    1      0      A/5 21171   7.2500   NaN      S      3
1  38.0    1      0      PC 17599  71.2833   C85      C      1
2  26.0    0      0  STON/O2. 3101282   7.9250   NaN      S      3
3  35.0    1      0      113803  53.1000  C123      S      1
4  35.0    0      0      373450   8.0500   NaN      S      3

      Survived
0      0
1      1
2      1
3      1
```

4 0

```
[31]: # size of the data
print("(row, column):",data.shape)
```

(row, column): (891, 12)

```
[32]: # check for missing values
data.isnull().sum()
```

```
[32]: PassengerId      0
      Name            0
      Sex            0
      Age           177
      SibSp          0
      Parch          0
      Ticket         0
      Fare           0
      Cabin          687
      Embarked       2
      Pclass         0
      Survived       0
      dtype: int64
```

```
[33]: # one option is to just simply drop rows with any missing values
data_no_missing_value=data.dropna()
print("(row, column) after dropping missing values:\n",data_no_missing_value.
      ↪shape)
```

(row, column) after dropping missing values:
(183, 12)

```
[34]: # get one column from the data
data_name=data_no_missing_value["Name"]
data_name
```

```
[34]: 1      Cumings, Mrs. John Bradley (Florence Briggs Th...
      3      Futrelle, Mrs. Jacques Heath (Lily May Peel)
      6      McCarthy, Mr. Timothy J
      10     Sandstrom, Miss. Marguerite Rut
      11     Bonnell, Miss. Elizabeth
      ...
      871    Beckwith, Mrs. Richard Leonard (Sallie Monypeny)
      872     Carlsson, Mr. Frans Olof
      879    Potter, Mrs. Thomas Jr (Lily Alexenia Wilson)
      887     Graham, Miss. Margaret Edith
      889     Behr, Mr. Karl Howell
```

Name: Name, Length: 183, dtype: object

```
[35]: # get three columns from the data
data_sex_age_survived=data_no_missing_value[['Sex', 'Age', 'Survived']]
data_sex_age_survived
```

```
[35]:
```

	Sex	Age	Survived
1	female	38.0	1
3	female	35.0	1
6	male	54.0	0
10	female	4.0	1
11	female	58.0	1
..
871	female	47.0	1
872	male	33.0	0
879	female	56.0	1
887	female	19.0	1
889	male	26.0	1

[183 rows x 3 columns]

```
[36]: # encode categorical attributes
data_sex_age_survived['Sex']=data_sex_age_survived.Sex.astype('category').cat.
      ↪codes
data_sex_age_survived
```

<ipython-input-36-9aac392383e8>:2: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.

Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
data_sex_age_survived['Sex']=data_sex_age_survived.Sex.astype('category').cat.
codes
```

```
[36]:
```

	Sex	Age	Survived
1	0	38.0	1
3	0	35.0	1
6	1	54.0	0
10	0	4.0	1
11	0	58.0	1
..
871	0	47.0	1
872	1	33.0	0
879	0	56.0	1
887	0	19.0	1
889	1	26.0	1

[183 rows x 3 columns]

```
[37]: # define X and y
feature_columns=['Sex', 'Age']
X=data_sex_age_survived[feature_columns]
y=data_sex_age_survived.Survived
print(X.head())
print(y.head())
```

	Sex	Age
1	0	38.0
3	0	35.0
6	1	54.0
10	0	4.0
11	0	58.0
1	1	
3	1	
6	0	
10	1	
11	1	

Name: Survived, dtype: int64

```
[38]: # train/test data split
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y)
```

```
[39]: # train a NN classifier
from sklearn import neighbors
clf = neighbors.KNeighborsClassifier(n_neighbors=1)
clf.fit(X_train, y_train)
```

```
[39]: KNeighborsClassifier(n_neighbors=1)
```

```
[40]: # predict the test data
y_pred=clf.predict(X_test)
y_pred
```

```
[40]: array([0, 1, 1, 0, 1, 1, 1, 0, 0, 0, 1, 1, 1, 1, 0, 1, 1, 1, 1, 0, 1, 1,
        0, 1, 1, 1, 1, 1, 0, 1, 0, 0, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,
        1, 1])
```

```
[41]: # calculate testing accuracy
from sklearn import metrics
print(metrics.accuracy_score(y_test, y_pred))
```

0.8043478260869565

[]:

[]:

2 Task 1: Build a k-NN classifier

```
[74]: # Load the dataset
data = pd.read_csv("/home/dalia/Master of Information Technology in_
↳Cybersecurity/Session 3/COMP8325 Artificial Intelligence/week_3/Week 3_
↳Workshop-20210310/data/titanic.csv")
print("data size: "+str(data.shape))
data.head()
```

data size: (891, 12)

```
[74]:
```

	PassengerId	Name	Sex	
0	1	Braund, Mr. Owen Harris	male	
1	2	Cumings, Mrs. John Bradley (Florence Briggs Th...	female	
2	3	Heikkinen, Miss. Laina	female	
3	4	Futrelle, Mrs. Jacques Heath (Lily May Peel)	female	
4	5	Allen, Mr. William Henry	male	

	Age	SibSp	Parch	Ticket	Fare	Cabin	Embarked	Pclass	
0	22.0	1	0	A/5 21171	7.2500	NaN	S	3	
1	38.0	1	0	PC 17599	71.2833	C85	C	1	
2	26.0	0	0	STON/O2. 3101282	7.9250	NaN	S	3	
3	35.0	1	0	113803	53.1000	C123	S	1	
4	35.0	0	0	373450	8.0500	NaN	S	3	

	Survived
0	0
1	1
2	1
3	1
4	0

```
[75]: # delete the attribute "Cabin"
del data['Cabin']
```

```
[76]: # fill missing values with the mean age
print("The mean age: ", data.Age.mean())

data.Age.fillna(data.Age.mean(), inplace=True)

#check for mmissing values
data.isnull().sum()
```

The mean age: 29.69911764705882

```
[76]: PassengerId    0
      Name          0
      Sex           0
      Age           0
      SibSp         0
      Parch         0
      Ticket        0
      Fare          0
      Embarked      2
      Pclass        0
      Survived      0
      dtype: int64
```

```
[77]: # drop rows with any missing values (the "Embarked" attribute)

data = data[data.Embarked.notnull()]

#check missing values
print(data.isnull().sum())

print("Data size with no missing values: "+str(data.shape))
```

```
PassengerId    0
Name           0
Sex            0
Age            0
SibSp          0
Parch          0
Ticket         0
Fare           0
Embarked       0
Pclass         0
Survived       0
dtype: int64
Data size with no missing values: (889, 11)
```

A question: Are “PassengerId”, “Name” and “Ticket” suitable for being a feature?

```
[78]: # Select all other attributes to form a data set
data_reduced = pd.DataFrame(data.drop(['PassengerId', 'Name', 'Ticket'],
→axis='columns'))
data_reduced.head()
```

```
[78]:      Sex  Age  SibSp  Parch    Fare  Embarked  Pclass  Survived
0  male  22.0     1     0   7.2500         S         3         0
1  female 38.0     1     0  71.2833         C         1         1
```

2	female	26.0	0	0	7.9250	S	3	1
3	female	35.0	1	0	53.1000	S	1	1
4	male	35.0	0	0	8.0500	S	3	0

```
[79]: # encode the categorical attributes where applicable
data_reduced['Sex'] = data_reduced['Sex'].astype('category').cat.codes
data_reduced['Embarked'] = data_reduced['Embarked'].astype('category').cat.codes
data_reduced.head()
```

```
[79]:
```

	Sex	Age	SibSp	Parch	Fare	Embarked	Pclass	Survived
0	1	22.0	1	0	7.2500	2	3	0
1	0	38.0	1	0	71.2833	0	1	1
2	0	26.0	0	0	7.9250	2	3	1
3	0	35.0	1	0	53.1000	2	1	1
4	1	35.0	0	0	8.0500	2	3	0

```
[80]: # build a 1-NN classification model where "Survived" is the prediction target
```

```
X = data_reduced.drop('Survived', axis='columns')
y = data_reduced['Survived']

# train and test data split
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y)

# build a 1-NN classification model

clf = neighbors.KNeighborsClassifier(n_neighbors = 1)
clf.fit(X_train, y_train)

y_pred=clf.predict(X_test)

print(metrics.accuracy_score(y_test, y_pred))
```

0.6995515695067265

```
[ ]:
```

3 Task 2 Tune the hyper parameters

```
[93]: # manually try other k values, e.g., 2, 3, 5, 10, the accuracy comparison
```

```
# build a 2-NN classification model
clf = neighbors.KNeighborsClassifier(n_neighbors = 2)
clf.fit(X_train, y_train)
```

```
y_pred=clf.predict(X_test)

print(metrics.accuracy_score(y_test, y_pred))
```

0.7130044843049327

```
[94]: # build a 3-NN classification model
      clf = neighbors.KNeighborsClassifier(n_neighbors = 3)
      clf.fit(X_train, y_train)

      y_pred=clf.predict(X_test)

      print(metrics.accuracy_score(y_test, y_pred))
```

0.7533632286995515

```
[95]: # build a 5-NN classification model
      clf = neighbors.KNeighborsClassifier(n_neighbors = 5)
      clf.fit(X_train, y_train)

      y_pred=clf.predict(X_test)

      print(metrics.accuracy_score(y_test, y_pred))
```

0.7354260089686099

```
[96]: # build a 10-NN classification model
      clf = neighbors.KNeighborsClassifier(n_neighbors = 10)
      clf.fit(X_train, y_train)

      y_pred=clf.predict(X_test)

      print(metrics.accuracy_score(y_test, y_pred))
      print(y_test)
      print(y_pred)
```

0.7174887892376681

```
141    1
169    0
95     0
726    1
242    0
..
609    1
206    0
668    0
461    0
400    1
```



```
Name: Survived, Length: 223, dtype: int64
[0 1 0 0 0 0 0 0 0 0 0 0 0 0 1 1 0 1 0 0 0 0 1 0 0 0 0 1 1 0 0 0 0 1 1 1 0
 0 0 0 0 0 0 0 0 0 0 1 0 0 0 1 0 0 0 0 0 0 0 0 0 0 0 1 0 1 1 0 0 1 0 1 0 0 0 0
 0 0 1 1 0 0 0 1 0 0 0 0 0 0 0 1 0 0 0 0 0 0 0 1 1 0 0 0 0 0 0 0 0 0 0 0 0 1 0 0
 0 1 0 1 0 0 0 0 0 1 0 0 0 1 0 0 1 0 1 1 0 0 0 0 0 0 0 1 0 1 0 0 0 0 0 0 0 1 0
 0 0 1 0 0 0 0 0 1 1 0 1 0 1 0 0 1 0 0 0 0 0 0 0 0 0 0 1 0 0 0 1 0 0 1 0 0 0
 1 1 0 0 0 1 0 1 1 0 0 0 0 0 0 0 0 0 1 0 1 0 1 1 0 0 0 0 0 0 0 0 0 0 1 0 0 0
 0]
```

```
[98]: # use k-fold cross-validation (10-fold cross validation)
from sklearn.model_selection import cross_val_score

n_values = range(1, 11)
for n in n_values:
    clf = neighbors.KNeighborsClassifier(n_neighbors = n)
    scores = cross_val_score(clf, X, y, cv=10, scoring='accuracy')
    print("Accuracy: %0.2f.(+/- %0.2f)" %(scores.mean(),scores.std()))
```


```
Accuracy: 0.69.(+/- 0.04)
Accuracy: 0.70.(+/- 0.04)
Accuracy: 0.72.(+/- 0.05)
Accuracy: 0.70.(+/- 0.04)
Accuracy: 0.70.(+/- 0.03)
Accuracy: 0.69.(+/- 0.03)
Accuracy: 0.72.(+/- 0.05)
Accuracy: 0.69.(+/- 0.04)
Accuracy: 0.71.(+/- 0.05)
Accuracy: 0.70.(+/- 0.04)
```

```
[86]: # visualize
import matplotlib.pyplot as plt

accuracy_scores = list()
accuracy_scores_std = list()

for n in n_values:
    clf = neighbors.KNeighborsClassifier(n_neighbors = n)
    scores = cross_val_score(clf, X, y, cv=10, scoring='accuracy')
    accuracy_scores.append(scores.mean())
    accuracy_scores_std.append(scores.std())

plt.errorbar(n_values, accuracy_scores, yerr=accuracy_scores_std, marker='o')
plt.title('Effect of neighbors')
plt.xlabel('n_neighbors')
plt.ylabel('Accuracy')
plt.ylim(0.6, 0.8)
plt.show()
```



```
[99]: # Use grid search to facilitate the identification of the best hyper parameter
from sklearn.model_selection import GridSearchCV

tuned_parameters = [{'n_neighbors': range(1,10)}]

search = GridSearchCV(neighbors.KNeighborsClassifier(), tuned_parameters,
    ↳scoring='accuracy', cv=10)
search.fit(X_train, y_train)

print("Best parameters set found on training set:")
print(search.best_params_)
```

```
Best parameters set found on training set:
{'n_neighbors': 8}
```

```
[ ]:
```

4 Task 3 Build a decision tree model for the above classification task

```
[ ]:
```

[]:

[]:

[]: